

1. Case study on Real-world Use Cases where ML model went badly wrong.
2. Perform Exploratory data analysis on structured, unstructured and image text data
3. Explore result visualization of post-hoc analysis methods:-
 - A. Partial dependence plot(PDP)
 - B. CNN:- layer-wise relevance propagation (LRP)(Guided backprop, Gradient CAM) Surrogate explainer
4. Explore result visualization of post-hoc analysis methods:-
 - A. Feature importance -sensitivity analysis
 - B. Counter factual examples
5. Implementing Data Centric XAI approach
6. Investigate the interpretability of LIME local explanation on tabular, image and text data
7. Explore different types of SHAP on ML and DL models to explore local and global explanations
8. Demonstrate the working of transformers using SHAP

3. Result Visualization (Post - Hoc Analysis Methods)

PDP

```
In [136... import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.inspection import PartialDependenceDisplay
import matplotlib.pyplot as plt
```

```
In [137... # Load California Housing Dataset
data = fetch_california_housing()
# Display Feature Names and Descriptions
print("Features in the California Housing Dataset:")
print(f"{'Feature Name':<15} Description")
print("="*50)
for name, desc in zip(data.feature_names, data.DESCR.split("\n")[12:20]):
    print(f"{'name':<15} {desc.strip()}")

X = pd.DataFrame(data.data, columns=data.feature_names)
y = data.target

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Train a Random Forest Regressor
model = RandomForestRegressor(random_state=42, n_estimators=100)
model.fit(X_train, y_train)
```

```
Features in the California Housing Dataset:
Feature Name      Description
=====
MedInc            - MedInc          median income in block group
HouseAge          - HouseAge        median house age in block group
AveRooms          - AveRooms        average number of rooms per household
AveBedrms         - AveBedrms       average number of bedrooms per household
Population        - Population      block group population
AveOccup          - AveOccup        average number of household members
Latitude          - Latitude        block group latitude
Longitude         - Longitude       block group longitude
```

```
Out[137... ▼ RandomForestRegressor ⓘ ⓘ
RandomForestRegressor(random_state=42)
```

```
In [138... # Generate Partial Dependence Plot for a single feature
plt.figure(figsize=(30, 30))
PartialDependenceDisplay.from_estimator(model, X_train, ['HouseAge'])
# plt.title("PDP for Feature: Median Income (MedInc)")
plt.show()
```


<Figure size 3000x3000 with 0 Axes>



```
In [140... # Generate 2D Partial Dependence Plot for two interacting features
plt.figure(figsize=(8, 6))
PartialDependenceDisplay.from_estimator(model, X_test, [("MedInc", "AveRooms")])
plt.title("2D PDP for Features: Median Income (MedInc) and Average Rooms (AveRooms)")
plt.show()
```

<Figure size 800x600 with 0 Axes>

2D PDP for Features: Median Income (MedInc) and Average Rooms (AveRooms)



CNN:- layer-wise relevance propagation (LRP)(Guided backprop, Gradient CAM) Surrogate explainer

```
In [67]: import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import mnist
import numpy as np
import matplotlib.pyplot as plt
```

```
In [133... # Load and normalize the data (with a validation split)
full_ds = tf.keras.utils.image_dataset_from_directory(
    'D:\SEM_6\XAI\LAB\image1',
    image_size=(28, 28),
    color_mode='grayscale',
    seed=123
).map(lambda x, y: (x / 255.0, y))

# Convert the dataset to numpy arrays
x_full, y_full = [], []
for images, labels in full_ds:
    x_full.append(images.numpy())
    y_full.append(labels.numpy())
x_full = np.concatenate(x_full, axis=0)
y_full = np.concatenate(y_full, axis=0)
```

```
# Split the data into training and testing sets (80% train, 20% test)
x_train, x_test, y_train, y_test = train_test_split(x_full, y_full, test_size=0.3, random_state=42)

# Optional: Print the shape of the data to verify
print(f"x_train shape: {x_train.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"x_test shape: {x_test.shape}")
print(f"y_test shape: {y_test.shape}")
```

Found 180 files belonging to 2 classes.

x_train shape: (126, 28, 28, 1)

y_train shape: (126,)

x_test shape: (54, 28, 28, 1)

y_test shape: (54,)

```
In [125... model = models.Sequential([
    layers.Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(32, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(10, activation='softmax')
])
```

```
In [127... model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.1)
```

Epoch 1/10

2/2 [=====] - 2s 431ms/step - loss: 0.7447 - accuracy: 0.5575 - val_loss: 0.8238 - val_accuracy: 0.5385

Epoch 2/10

2/2 [=====] - 0s 108ms/step - loss: 0.6428 - accuracy: 0.6106 - val_loss: 1.1020 - val_accuracy: 0.4615

Epoch 3/10

2/2 [=====] - 0s 106ms/step - loss: 0.6800 - accuracy: 0.6106 - val_loss: 0.9205 - val_accuracy: 0.4615

Epoch 4/10

2/2 [=====] - 0s 110ms/step - loss: 0.6037 - accuracy: 0.6283 - val_loss: 0.7992 - val_accuracy: 0.5385

Epoch 5/10

2/2 [=====] - 0s 129ms/step - loss: 0.5786 - accuracy: 0.7168 - val_loss: 0.7911 - val_accuracy: 0.5385

Epoch 6/10

2/2 [=====] - 0s 141ms/step - loss: 0.5884 - accuracy: 0.7080 - val_loss: 0.7892 - val_accuracy: 0.5385

Epoch 7/10

2/2 [=====] - 0s 143ms/step - loss: 0.5687 - accuracy: 0.7168 - val_loss: 0.8112 - val_accuracy: 0.4615

Epoch 8/10

2/2 [=====] - 0s 206ms/step - loss: 0.5601 - accuracy: 0.6637 - val_loss: 0.8487 - val_accuracy: 0.3846

Epoch 9/10

2/2 [=====] - 0s 152ms/step - loss: 0.5729 - accuracy: 0.6726 - val_loss: 0.8638 - val_accuracy: 0.4615

Epoch 10/10

2/2 [=====] - 0s 144ms/step - loss: 0.5731 - accuracy: 0.6637 - val_loss: 0.8340 - val_accuracy: 0.3077

Out[127... <keras.src.callbacks.History at 0x1fab592a400>

Grad-CAM

```
In [128... from tf_keras_vis.gradcam import Gradcam
from tf_keras_vis.utils.model_modifiers import ReplaceToLinear
from tf_keras_vis.utils.scores import CategoricalScore
```


```
In [129... # Define the GradCAM object
# replace2linear = ReplaceToLinear()
gradcam = Gradcam(model, model_modifier=ReplaceToLinear())

# Select test image
image = x_test[0:1]
label = y_test[0]
score = CategoricalScore([label])

# Generate heatmap
cam = gradcam(score, image) # auto-detects last conv layer
heatmap = cam[0]

# Plot result
plt.imshow(image[0], cmap='gray')
plt.imshow(heatmap, alpha=0.5)
plt.title("Grad-CAM")
plt.axis('off')
plt.show()
```

Grad-CAM



Guided Back Prop

```
In [130... from tf_keras_vis.utils.model_modifiers import GuidedBackpropagation
from tf_keras_vis.saliency import Saliency
from tf_keras_vis.utils.scores import CategoricalScore
```


```
In [131... # Guided Backpropagation setup

saliency = Saliency(model, model_modifier=GuidedBackpropagation())

# Generate saliency map (Guided Backprop)
image = x_test[0:1]
score = CategoricalScore([y_test[0]])
gbp = saliency(score, image)

# Visualize
plt.imshow(np.abs(gbp[0].squeeze()), cmap='viridis')
plt.title("Guided Backpropagation")
plt.axis('off')
plt.show()
```

Guided Backpropagation



Layer-Wise relevance propagation

```
In [132... img = x_test[0] # shape: (28, 28, 1)
img_batch = np.expand_dims(img, axis=0)

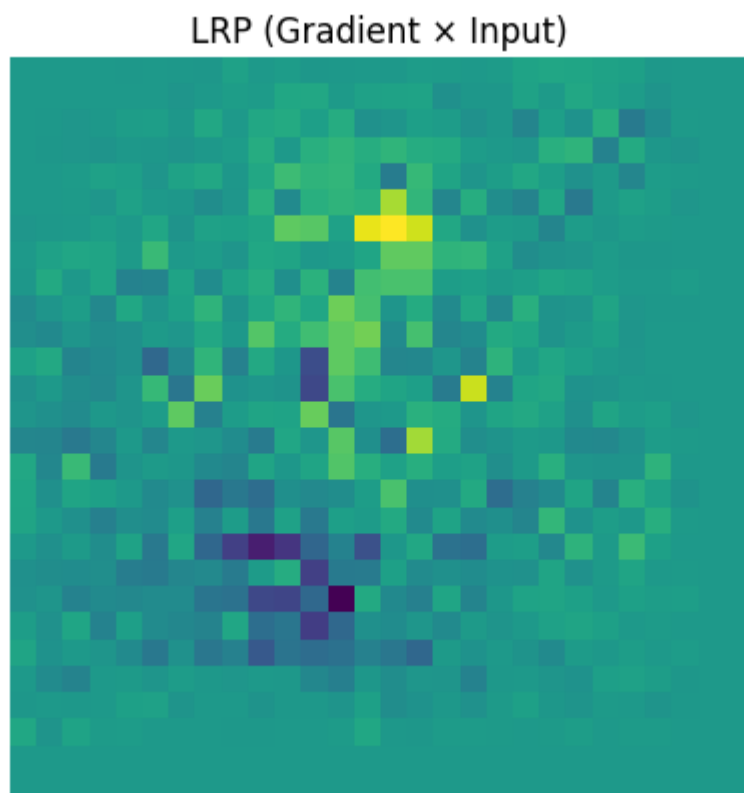
# Gradient x Input (Basic LRP approximation)
img_tensor = tf.convert_to_tensor(img_batch)

with tf.GradientTape() as tape:
    tape.watch(img_tensor)
    preds = model(img_tensor)
    class_idx = tf.argmax(preds[0])
    loss = preds[0, class_idx]
```

```
# Compute gradients
grads = tape.gradient(loss, img_tensor)[0].numpy() # shape: (28, 28, 1)

relevance = grads * img # same shape

# Plot result
plt.imshow(relevance)
plt.title("LRP (Gradient x Input)")
plt.axis('off')
plt.show()
```



4. Explore result visualization of post-hoc analysis methods

Feature importance -sensitivity analysis

```
In [98]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
```

```
In [105... # Load dataset
data = load_iris(as_frame=True)
df = data.frame
df['target'] = data.target


# Use only 2 classes for simplicity
df = df[df['target'].isin([0, 1]).reset_index(drop=True)

# Split data
X = df.drop('target', axis=1)
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
model = RandomForestClassifier().fit(X_train, y_train)
```

```
In [106... feat_imp=model.feature_importances_
feat=X.columns
plt.barh(feat,feat_imp)
```

```
Out[106... <BarContainer object of 4 artists>
```



Counterfactual examples

```
In [89]: from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import pandas as pd
```

```
In [90]: # Prepare dataset
data = load_iris(as_frame=True)
df = data.frame
df['target'] = data.target

# Use only 2 classes for simplicity
df = df[df['target'].isin([0, 1]).reset_index(drop=True)

# Split data
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

# Train model
model = RandomForestClassifier()
model.fit(X_train, y_train)
```

```
Out[90]: ▼ RandomForestClassifier ⓘ ?
RandomForestClassifier()
```

```
In [91]: import dice_ml
from dice_ml.utils import helpers
```

```
In [88]: # Wrap in DiCE
d = dice_ml.Data(dataframe=df, continuous_features=X.columns.tolist(), outcome_name='target')
m = dice_ml.Model(model=model, backend='sklearn')

# Generate counterfactuals
explainer = dice_ml.Dice(d, m, method='random')
query_instance = pd.DataFrame([X_test.iloc[0]])

cf = explainer.generate_counterfactuals(query_instance, total_CFs=2, desired_class="opposite")
cf.visualize_as_dataframe(show_only_changes=True)
```

100% | 1/1 [00:00<00:00, 2.43it/s]

Query instance (original outcome : 0)

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.0	3.4	1.6	0.4	0

Diverse Counterfactual set (new outcome: 1)

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	-	3.6	-	1.7	1.0
1	-	-	5.1	1.6	1.0

5. Implementing Data Centric XAI approach

In [107... `## NO DeepChecks So, no Data Centric For Model LAB`

6. Investigate the interpretability of LIME local explanation on tabular, image and text data

Tabular

In [144... `import pandas as pd
import numpy as np
import lime
import lime.lime_tabular
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder

Load dataset (download from Kaggle or GitHub and place Locally)
df = pd.read_csv('Loan_Pred.csv') # External dataset

Simple preprocessing
df.dropna(inplace=True)
label_cols = ['Gender', 'Married', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']
for col in label_cols:
 df[col] = LabelEncoder().fit_transform(df[col])
df.describe()`

Out[144...

	Gender	Married	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	480.000000	480.000000	480.000000	480.000000	480.000000	480.000000	480.000000	480.000000	480.000000
mean	0.820833	0.647917	0.202083	0.137500	5364.231250	1581.093583	144.735417	342.050000	0.009792
std	0.383892	0.478118	0.401973	0.344734	5668.251251	2617.692267	80.508164	65.212401	0.342021
min	0.000000	0.000000	0.000000	0.000000	150.000000	0.000000	9.000000	36.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	2898.750000	0.000000	100.000000	360.000000	0.000000
50%	1.000000	1.000000	0.000000	0.000000	3859.000000	1084.500000	128.000000	360.000000	0.000000
75%	1.000000	1.000000	0.000000	0.000000	5852.500000	2253.250000	170.000000	360.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	81000.000000	33837.000000	600.000000	480.000000	1.000000

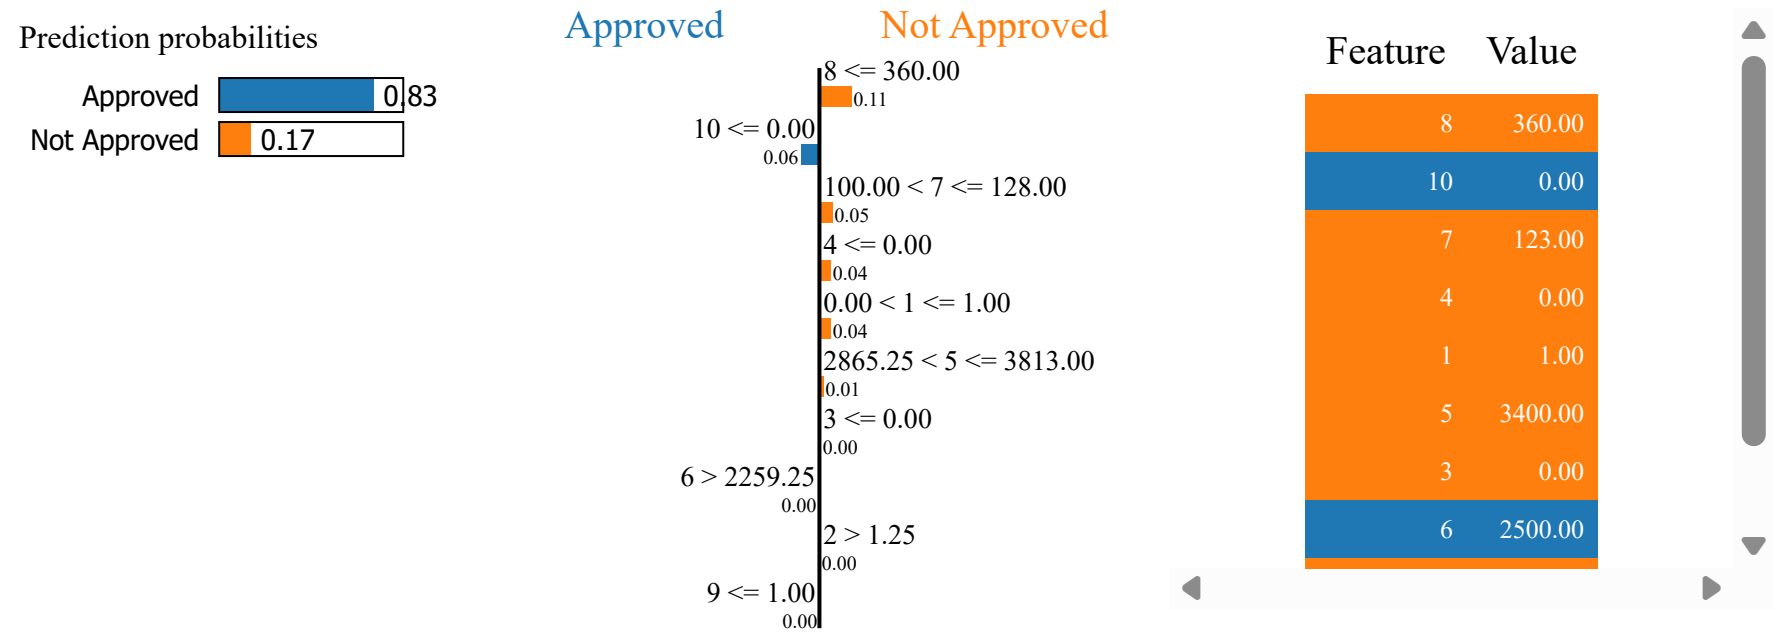
In [145... `df.drop('Loan_ID',axis=1,inplace=True)
df=df.replace('3+',3).astype(int)
X = df.drop(['Loan_Status'], axis=1)
y = df['Loan_Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

model = RandomForestClassifier().fit(X_train, y_train)`

In [146... `lime_exp=lime.lime_tabular.LimeTabularExplainer(
 training_data=X_train.values,
 training_labels=X_train.columns,
 mode='classification',
 class_names=['Approved', 'Not Approved']
)# LIME Explanation

exp=lime_exp.explain_instance(X_test.iloc[3].values,model.predict_proba)
exp.show_in_notebook()`

C:\Users\saiha\anaconda3\envs\XAI_Env\lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names
warnings.warn(



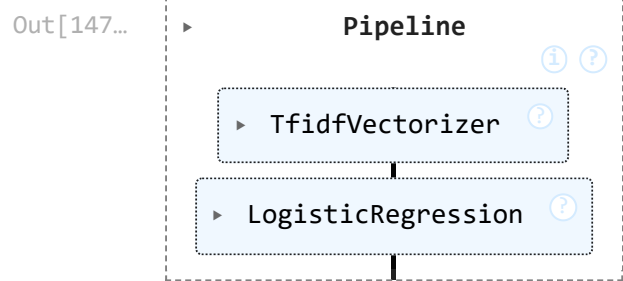
Text data

```
In [147... import pandas as pd
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression

# Load dataset
df = pd.read_csv('sms.tsv', sep='\t', header=None, names=['label', 'message'])

df['label'] = df['label'].map({'ham': 0, 'spam': 1})
X_train, X_test, y_train, y_test = train_test_split(df['message'], df['label'], test_size=0.2)

# Pipeline
vectorizer = TfidfVectorizer()
model = LogisticRegression()
pipe = make_pipeline(vectorizer, model)
pipe.fit(X_train, y_train)
```



```
In [148... import lime.lime_text
# LIME Explanation
explainer = lime.lime_text.LimeTextExplainer(class_names=['Ham', 'Spam'])
exp = explainer.explain_instance(X_train.iloc[5], pipe.predict_proba)
exp.show_in_notebook()
```




Image Data

```
In [153... from lime import lime_image
from skimage.segmentation import mark_boundaries
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.mobilenet import preprocess_input
from tensorflow.keras.models import load_model
import matplotlib.pyplot as plt
import numpy as np
```



```

model = load_model('mobilenet_model.h5') # Assume trained on cats vs dogs
img_path = 'lion.jpg'
img = image.load_img(img_path, target_size=(224, 224))
img_array = image.img_to_array(img)
img_preprocessed = preprocess_input(np.expand_dims(img_array, axis=0))

```

WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.

```

In [169... explainer = lime_image.LimeImageExplainer()
explanation = explainer.explain_instance(
    image=img_array,
    classifier_fn=model.predict,
    num_samples=10
)

# Show explanation for top label
temp,mask = explanation.get_image_and_mask(
    label=explanation.top_labels[0]
)

plt.imshow(mark_boundaries(temp / 255.0, mask))
plt.title('LIME Explanation')
plt.axis('off')
plt.show()


```

100%  10/10 [00:00<00:00, 18.01it/s]

1/1 [=====] - 0s 407ms/step

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003 921569..1.0].

LIME Explanation



```

In [167... def draw_boundaries(image, mask):
    boundaries = ndimage.binary_dilation(mask) ^ mask # Edges of mask
    outlined_image = image.copy()

    # Draw red boundary on image
    outlined_image[boundaries] = [255, 0, 0] # Red outline


    return outlined_image

highlighted = draw_boundaries(temp, mask)

plt.imshow(highlighted.astype(np.uint8))
plt.title('LIME Explanation (Custom Boundary)')
plt.axis('off')
plt.show()

```

LIME Explanation (Custom Boundary)



8. Demonstrate the working of transformers using SHAP

```
In [170... from transformers import pipeline, AutoModelForSequenceClassification, AutoTokenizer
import shap

model_name = "distilbert-base-uncased-finetuned-sst-2-english"
classifier = pipeline("sentiment-analysis", model=model_name)
explainer = shap.Explainer(classifier)

shap_values = explainer(["I love using SHAP with transformers"])
shap.plots.text(shap_values)
```

Device set to use cpu

